

Challenges Faced

- Configuring Hadoop services and resolving the SSH connection issue during Hadoop startup.
- Handling the “**Output directory already exists**” error while executing the Hadoop Streaming job.
- Managing HDFS input and output directories for different dataset sizes.
- Ensuring the Python Mapper and Reducer scripts had the correct permissions and executed successfully.
- Measuring and comparing the actual MapReduce processing time for 1 MB, 10 MB, and 100 MB datasets.