

Practical 7

Problem definition:

You are a Big Data Engineer at MovieFlix Technologies, a movie recommendation company serving millions of users worldwide. The recommendation engine repeatedly processes movie ratings data to generate personalized suggestions. Due to the increasing volume of user interactions, query execution time has become a major concern. Management wants to optimize the Spark application using Caching and Partitioning techniques to reduce execution time and improve scalability.

Objectives:

1. To create a Spark Session and load the movie ratings dataset into a Spark DataFrame.
2. To display and understand the schema of the dataset.
3. To check the initial number of partitions.
4. To calculate the average movie ratings.
5. To measure execution time before optimization.
6. To apply cache() and trigger caching using count().
7. To measure execution time after caching.
8. To repartition the dataset and observe the change in the number of partitions.
9. To compare the performance before and after caching and partitioning.
10. To understand the effect of cache(), persist(), repartition(), and coalesce() on Spark performance.

Observation / Output:

Task 1: Create Spark Session.

CODE :

```
from pyspark.sql import SparkSession

# Create Spark Session
spark = SparkSession.builder \
    .appName("MovieFlix_Caching_Partitioning") \
    .getOrCreate()

print("=====")
print("    MovieFlix Technologies")
print("=====")
print("Spark Session Created Successfully!")
print("Spark Version:", spark.version)
print("=====")

# Stop Spark Session
spark.stop()
```

```

26/09/29 18:37:03 INFO DiskBlockManager: Created local directory at /tmp/blockmgr-f62942ed-e749-443b-b748-08da803a66dd
26/09/29 18:37:03 INFO MemoryStore: MemoryStore started with capacity 434.4 MiB
26/09/29 18:37:03 INFO SparkEnv: Registering OutputCommitCoordinator
26/09/29 18:37:03 INFO JettyUtils: Start Jetty 0.0.0.0:4040 for SparkUI
26/09/29 18:37:03 INFO Utils: Successfully started service 'SparkUI' on port 4040.
26/09/29 18:37:03 INFO Executor: Starting executor ID driver on host 10.255.255.254
26/09/29 18:37:03 INFO Executor: OS info Linux, 6.18.33.2-microsoft-standard-WSL2, amd64
26/09/29 18:37:03 INFO Executor: Java version 21.0.11
26/09/29 18:37:03 INFO Executor: Starting executor with user classpath (userClassPathFirst = false): ''
26/09/29 18:37:03 INFO Executor: Created or updated repl class loader org.apache.spark.util.MutableURLClassLoader@224a7536 for default.
26/09/29 18:37:03 INFO NettyBlockTransferService: Server created on 10.255.255.254:39109
26/09/29 18:37:03 INFO BlockManager: Using org.apache.spark.storage.RandomBlockReplicationPolicy for block replication policy
26/09/29 18:37:03 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, 10.255.255.254, 39109, None)
26/09/29 18:37:03 INFO BlockManagerMasterEndpoint: Registering block manager 10.255.255.254:39109 with 434.4 MiB RAM, BlockManagerId(driver, 10.255.255.254, 39109, None)
26/09/29 18:37:03 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, 10.255.255.254, 39109, None)
26/09/29 18:37:03 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, 10.255.255.254, 39109, None)

MovieFlix Technologies

Spark Session Created Successfully!
Spark Version: 3.5.7

26/09/29 18:37:04 INFO SparkContext: SparkContext is stopping with exitCode 0.
26/09/29 18:37:04 INFO SparkUI: Stopped Spark web UI at http://10.255.255.254:4040
26/09/29 18:37:04 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
26/09/29 18:37:04 INFO MemoryStore: MemoryStore cleared
26/09/29 18:37:04 INFO BlockManager: BlockManager stopped
26/09/29 18:37:04 INFO BlockManagerMaster: BlockManagerMaster stopped
26/09/29 18:37:04 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
26/09/29 18:37:04 INFO SparkContext: Successfully stopped SparkContext
26/09/29 18:37:04 INFO ShutdownHookManager: Shutdown hook called
26/09/29 18:37:04 INFO ShutdownHookManager: Deleting directory /tmp/spark-8833ddeb-3d22-4f8a-ad96-39d38861823a/pyspark-fd92b985-8c94-4e43-a37d-bc
26/09/29 18:37:04 INFO ShutdownHookManager: Deleting directory /tmp/spark-8833ddeb-3d22-4f8a-ad96-39d38861823a
26/09/29 18:37:04 INFO ShutdownHookManager: Deleting directory /tmp/spark-b3612e78-e7a0-4d5a-a997-786a0ee22127

```

Task 2: Load ratings dataset into Spark DataFrame. Display dataset schema.
from pyspark.sql import SparkSession

```

spark = SparkSession.builder \
    .appName("MovieFlix_Load_Ratings") \
    .getOrCreate()

print("    MovieFlix Technologies")

ratings_df = spark.read.csv(
    "movie ratings.csv",
    header=True,
    inferSchema=True
)

print("\nRatings Dataset Loaded Successfully!")

print("\n===== FIRST 5 RECORDS =====")

ratings_df.show(5)

print("\n===== DATASET SCHEMA =====")

ratings_df.printSchema()

spark.stop()

```

```

26/09/29 21:08:56 INFO TaskSchedulerImpl: Killing all running tasks in stage 2: Stage finished
26/09/29 21:08:56 INFO DAGScheduler: Job 2 finished: showString at NativeMethodAccessorImpl.java:0, took 0.121981 s
26/09/29 21:08:56 INFO CodeGenerator: Code generated in 17.636119 ms
+-----+-----+-----+-----+
|userId|movieId|rating|timestamp|
+-----+-----+-----+-----+
|1|16|4.0|1217897793|
|1|24|1.5|1217895807|
|1|32|4.0|1217896246|
|1|47|4.0|1217896556|
|1|50|4.0|1217896523|
+-----+-----+-----+-----+
only showing top 5 rows

PROJECT: DATASET SCHEMA
root
 |-- userId: integer (nullable = true)
 |-- movieId: integer (nullable = true)
 |-- rating: double (nullable = true)
 |-- timestamp: integer (nullable = true)

26/09/29 21:08:56 INFO SparkContext: SparkContext is stopping with exitCode 0.
26/09/29 21:08:56 INFO SparkUI: Stopped Spark web UI at http://10.255.255.254:4040
26/09/29 21:08:56 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
26/09/29 21:08:56 INFO MemoryStore: MemoryStore cleared
26/09/29 21:08:56 INFO BlockManager: BlockManager stopped
26/09/29 21:08:56 INFO BlockManagerMaster: BlockManagerMaster stopped
26/09/29 21:08:56 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
26/09/29 21:08:56 INFO SparkContext: Successfully stopped SparkContext
26/09/29 21:08:56 INFO ShutdownHookManager: Shutdown hook called
26/09/29 21:08:56 INFO ShutdownHookManager: Deleting directory /tmp/spark-91aa5ccf-0526-406b-8685-dd7abd59f22d
26/09/29 21:08:56 INFO ShutdownHookManager: Deleting directory /tmp/spark-e62c2f5a-c969-40bf-8e7e-e13e79c6b22a
26/09/29 21:08:56 INFO ShutdownHookManager: Deleting directory /tmp/spark-e62c2f5a-c969-40bf-8e7e-e13e79c6b22a/pyspark-770a8a9b-59cc-4576-b881-4d712f618e08
(tushar@Tushar-laptop)-[~/Desktop/BDA_Practical7]

```

Task 3: Check Current Number of Partitions

from pyspark.sql import SparkSession

Create Spark Session

```

spark = SparkSession.builder \
    .appName("MovieFlix_Check_Partitions") \
    .getOrCreate()

```

```
print("    MovieFlix Technologies")
```

Load ratings dataset

```

ratings_df = spark.read.csv(
    "movie ratings.csv",
    header=True,
    inferSchema=True
)

```

Check current number of partitions

```
initial_partitions = ratings_df.rdd.getNumPartitions()
```

```
print("\n===== CURRENT PARTITIONS =====")
```

```
print("Number of partitions:", initial_partitions)
```

Stop Spark

```
spark.stop()
```

```

26/09/29 21:17:46 INFO FileScanRDD: Reading File path: file:///home/tushar/Desktop/BDA_Practical7/movie%20ratings.csv, range: 0-2475052, partition values: [empty row]
26/09/29 21:17:46 INFO BlockManagerInfo: Removed broadcast_1_piece0 on 10.255.255.254:40509 in memory (size: 6.4 KiB, free: 434.3 MiB)
26/09/29 21:17:46 INFO BlockManagerInfo: Removed broadcast_0_piece0 on 10.255.255.254:40509 in memory (size: 34.3 KiB, free: 434.4 MiB)
26/09/29 21:17:46 INFO Executor: Finished task 0.0 in stage 1.0 (TID 1). 1664 bytes result sent to driver
26/09/29 21:17:46 INFO TaskSetManager: Finished task 0.0 in stage 1.0 (TID 1) in 308 ms on 10.255.255.254 (executor driver) (1/1)
26/09/29 21:17:46 INFO TaskSchedulerImpl: Removed TaskSet 1.0, whose tasks have all completed, from pool
26/09/29 21:17:46 INFO DAGScheduler: ResultStage 1 (csv at NativeMethodAccessorImpl.java:0) finished in 0.350 s
26/09/29 21:17:46 INFO DAGScheduler: Job 1 is finished. Cancelling potential speculative or zombie tasks for this job
26/09/29 21:17:46 INFO TaskSchedulerImpl: Killing all running tasks in stage 1: Stage finished
26/09/29 21:17:46 INFO DAGScheduler: Job 1 finished: csv at NativeMethodAccessorImpl.java:0, took 0.357943 s
26/09/29 21:17:46 INFO FileSourceStrategy: Pushed Filters:
26/09/29 21:17:46 INFO FileSourceStrategy: Post-Scan Filters:
26/09/29 21:17:46 INFO MemoryStore: Block broadcast_4 stored as values in memory (estimated size 199.5 KiB, free 433.9 MiB)
26/09/29 21:17:46 INFO MemoryStore: Block broadcast_4_piece0 stored as bytes in memory (estimated size 34.3 KiB, free 433.9 MiB)
26/09/29 21:17:46 INFO BlockManagerInfo: Added broadcast_4_piece0 in memory on 10.255.255.254:40509 (size: 34.3 KiB, free: 434.3 MiB)
26/09/29 21:17:46 INFO SparkContext: Created broadcast 4 from javaToPython at NativeMethodAccessorImpl.java:0
26/09/29 21:17:46 INFO FileSourceScanExec: Planning scan with bin packing, max size: 4194304 bytes, open cost is considered as scanning 4194304 bytes.

===== CURRENT PARTITIONS =====
Number of partitions: 1
26/09/29 21:17:46 INFO SparkContext: SparkContext is stopping with exitCode 0.
26/09/29 21:17:46 INFO SparkUI: Stopped Spark web UI at http://10.255.255.254:4040
26/09/29 21:17:46 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
26/09/29 21:17:46 INFO MemoryStore: MemoryStore cleared
26/09/29 21:17:46 INFO BlockManager: BlockManager stopped
26/09/29 21:17:46 INFO BlockManagerMaster: BlockManagerMaster stopped
26/09/29 21:17:46 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
26/09/29 21:17:46 INFO SparkContext: Successfully stopped SparkContext
26/09/29 21:17:46 INFO ShutdownHookManager: Shutdown hook called
26/09/29 21:17:46 INFO ShutdownHookManager: Deleting directory /tmp/spark-b32a9219-ff2b-4322-8cf2-198c33a1fd9/pyspark-822e3a33-a99f-4821-a76e-18c0fb8df398
26/09/29 21:17:46 INFO ShutdownHookManager: Deleting directory /tmp/spark-9e5616b8-632c-4cca-af58-f20486331c08
26/09/29 21:17:46 INFO ShutdownHookManager: Deleting directory /tmp/spark-b32a9219-ff2b-4322-8cf2-198c33a1fd9

tushar@Tushar-laptop: [~/Desktop/BDA_Practical7]

```

Task 4: Calculate Average Movie Ratings

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import avg

```

```
# Create Spark Session
```

```

spark = SparkSession.builder \
    .appName("MovieFlix_Average_Ratings") \
    .getOrCreate()

```

```
print("    MovieFlix Technologies")
```

```
# Load ratings dataset
```

```

ratings_df = spark.read.csv(
    "movie ratings.csv",
    header=True,
    inferSchema=True
)

```

```
# Calculate average rating for each movie
```

```

average_ratings = ratings_df.groupBy("movieId") \
    .agg(avg("rating").alias("average_rating"))

```

```
print("\n AVERAGE MOVIE RATINGS ")
```

```
average_ratings.show()
```

```
# Stop Spark
```

```
spark.stop()
```

```

26/09/29 21:19:52 INFO DAGScheduler: Job 3 is finished. Cancelling potential speculative or zombie tasks for this job
26/09/29 21:19:52 INFO TaskSchedulerImpl: Killing all running tasks in stage 4: Stage finished
26/09/29 21:19:52 INFO DAGScheduler: Job 3 finished: showString at NativeMethodAccessorImpl.java:0, took 0.174992 s
26/09/29 21:19:52 INFO CodeGenerator: Code generated in 9.888819 ms

```

movieId	average_rating
1580	3.627906976744186
471	3.6666666666666665
7253	4.0
3997	2.5
32460	4.5
44022	3.3125
54190	3.357142857142857
833	2.95
1591	2.257575757575758
6357	3.6
2366	3.532258064516129
1088	3.419642857142857
1238	3.75
4519	3.0454545454545454
1645	3.5666666666666667
8638	3.9722222222222223
3175	3.712328767123288
2122	2.2142857142857144
26755	0.5
5803	2.5714285714285716

```

only showing top 20 rows

```

```

26/09/29 21:19:52 INFO SparkContext: SparkContext is stopping with exitCode 0.
26/09/29 21:19:52 INFO BlockManagerInfo: Removed broadcast_5_piece0 on 10.255.255.254:46779 in memory (size: 19.6 KiB, free: 434.3 MiB)
26/09/29 21:19:52 INFO BlockManagerInfo: Removed broadcast_6_piece0 on 10.255.255.254:46779 in memory (size: 21.4 KiB, free: 434.4 MiB)
26/09/29 21:19:53 INFO SparkUI: Stopped Spark web UI at http://10.255.255.254:4040
26/09/29 21:19:53 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
26/09/29 21:19:53 INFO MemoryStore: MemoryStore cleared
26/09/29 21:19:53 INFO BlockManager: BlockManager stopped
26/09/29 21:19:53 INFO BlockManagerMaster: BlockManagerMaster stopped

```

Task 5: Measure Execution Time Before Optimization.

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import avg

```

```

spark = SparkSession.builder \
    .appName("MovieFlix_Caching") \
    .getOrCreate()
print("    MovieFlix Technologies")

```

```

ratings_df = spark.read.csv(
    "movie ratings.csv",
    header=True,
    inferSchema=True
)

```

```

print("\nRatings Dataset Loaded Successfully!")

```

```

print("\n===== APPLYING CACHE =====")

```

```

ratings_df.cache()
print("cache() applied successfully.")
record_count = ratings_df.count()
print("Caching triggered using count().")
print("Total Records:", record_count)
print("\n===== AVERAGE RATINGS AFTER CACHING =====")
average_ratings = ratings_df.groupBy("movieId") \
    .agg(avg("rating").alias("average_rating"))

```

```

average_ratings.show()

```

```

spark.stop()

```

The screenshot shows a Jupyter Notebook with a DataFrame containing movie ratings. The top 20 rows are displayed, showing columns for movie ID, title, and rating. Below the DataFrame, the execution time before optimization is shown as 1.457803 seconds. The logs indicate that the SparkContext, SparkUI, MapOutputTrackerMasterEndpoint, MemoryStore, BlockManager, BlockManagerMaster, OutputCommitCoordinator, and ShutdownHookManager have been successfully stopped or deleted.

```

Session Actions Edit View Help
54190 3.3571428571428571
833 2.95
1591 2.257575757575758
6357 3.6
2366 3.532258064516129
1088 3.4196428571428571
1238 3.75
4519 3.0454545454545454
1645 3.566666666666667
8638 3.9722222222222223
3175 3.71232876123288
2122 2.2142857142857144
26755 0.5
5803 2.5714285714285716
+-----+
only showing top 20 rows

EXECUTION TIME BEFORE OPTIMIZATION
Execution Time BEFORE Optimization: 1.457803 seconds
26/09/29 21:22:33 INFO SparkContext: SparkContext is stopping with exitCode 0.
26/09/29 21:22:33 INFO SparkUI: Stopped Spark web UI at http://10.255.255.254:4040
26/09/29 21:22:33 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
26/09/29 21:22:33 INFO MemoryStore: MemoryStore cleared
26/09/29 21:22:33 INFO BlockManager: BlockManager stopped
26/09/29 21:22:33 INFO BlockManagerMaster: BlockManagerMaster stopped
26/09/29 21:22:33 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
26/09/29 21:22:33 INFO SparkContext: Successfully stopped SparkContext
26/09/29 21:22:33 INFO ShutdownHookManager: Shutdown hook called
26/09/29 21:22:33 INFO ShutdownHookManager: Deleting directory /tmp/spark-08c31eba-7d37-4cfb-8264-618b10fa3c3d
26/09/29 21:22:33 INFO ShutdownHookManager: Deleting directory /tmp/spark-7309365d-e4c7-4220-a628-e66789742ce3
26/09/29 21:22:33 INFO ShutdownHookManager: Deleting directory /tmp/spark-7309365d-e4c7-4220-a628-e66789742ce3/pyspark-d5847af0-1b40-4a5e-a30d-e0e6f8040922
tushar@Tushar-laptop: ~/Desktop/BDA_Practical7

```

Task 6: Apply cache() on the DataFrame. Trigger caching using count(). Recalculate average ratings.

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import avg
```

```
spark = SparkSession.builder \
    .appName("MovieFlix_Caching") \
    .getOrCreate()
```

```
# Show only ERROR logs
spark.sparkContext.setLogLevel("ERROR")
```

```
print("    MovieFlix Technologies")
```

```
ratings_df = spark.read.csv(
    "movie ratings.csv",
    header=True,
    inferSchema=True
)
```

```
print("\nRatings Dataset Loaded Successfully!")
```

```
print("\n===== APPLYING CACHE =====")
```

```
ratings_df.cache()
```

```
print("cache() applied successfully.")
record_count = ratings_df.count()
```


```
print("Caching triggered using count().")
print("Total Records:", record_count)
```

```
print("\n===== AVERAGE RATINGS AFTER CACHING =====")
```

```
average_ratings = ratings_df.groupBy("movieId") \
    .agg(avg("rating").alias("average_rating"))
```

```
average_ratings.show()
```

```
spark.stop()
```



Task 7: Measure execution time after caching. Repartition the dataset. Display updated partition count. Compare performance before and after optimization. Generate performance report.

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import avg
import time

spark = SparkSession.builder \
    .appName("MovieFlix_Performance_Optimization") \
    .getOrCreate()

# Show only errors, hide Spark INFO logs
spark.sparkContext.setLogLevel("ERROR")

print("    MOVIEFLIX TECHNOLOGIES")
print("    SPARK PERFORMANCE OPTIMIZATION")

ratings_df = spark.read.csv(
    "movie ratings.csv",
    header=True,
    inferSchema=True
)

print("\nDataset Loaded Successfully!")
```

```
initial_partitions = ratings_df.rdd.getNumPartitions()

print("\n-----")
print("INITIAL PARTITION COUNT")
print("-----")
print("Initial Partitions:", initial_partitions)

print("\n-----")
print("EXECUTION BEFORE OPTIMIZATION")
print("-----")

start_time = time.perf_counter()

average_before = ratings_df.groupBy("movieId") \
    .agg(avg("rating").alias("average_rating"))

# Trigger execution
average_before.count()

end_time = time.perf_counter()

before_optimization = end_time - start_time

print("Average ratings calculated successfully.")
print(f"Execution Time BEFORE Optimization: "
      f"{before_optimization:.6f} seconds")

print("\n-----")
print("APPLYING CACHE")
print("-----")

ratings_df.cache()

print("cache() applied successfully.")

cache_start = time.perf_counter()

total_records = ratings_df.count()

cache_end = time.perf_counter()

cache_time = cache_end - cache_start

print("Caching triggered using count().")
print("Total Records:", total_records)
print(f"Cache Time: {cache_time:.6f} seconds")

print("\n-----")
print("AVERAGE RATINGS AFTER CACHING")
print("-----")
```

```
start_time = time.perf_counter()

average_after_cache = ratings_df.groupBy("movieId") \
    .agg(avg("rating").alias("average_rating"))

average_after_cache.count()

end_time = time.perf_counter()

after_caching = end_time - start_time

average_after_cache.show(10, truncate=False)

print(f"Execution Time AFTER Caching: "
      f"{after_caching:.6f} seconds")

print("\n-----")
print("REPARTITIONING DATASET")
print("-----")

repartitioned_df = ratings_df.repartition(4)

updated_partitions = repartitioned_df.rdd.getNumPartitions()

print("Dataset repartitioned successfully.")
print("Updated Partition Count:", updated_partitions)

print("\n-----")
print("AVERAGE RATINGS AFTER PARTITIONING")
print("-----")

start_time = time.perf_counter()

average_after_partition = repartitioned_df.groupBy("movieId") \
    .agg(avg("rating").alias("average_rating"))

average_after_partition.count()

end_time = time.perf_counter()

after_partitioning = end_time - start_time

average_after_partition.show(10, truncate=False)

print(f"Execution Time AFTER Partitioning: "
      f"{after_partitioning:.6f} seconds")

cache_improvement = (
    (before_optimization - after_caching)
    / before_optimization
```

```
) * 100

partition_improvement = (
    (before_optimization - after_partitioning)
    / before_optimization
) * 100

print("\n")
print("          PERFORMANCE REPORT")
print("Dataset          : Movie Ratings Dataset")
print("Optimization Techniques : Caching and Partitioning")

print("\nPartition Information")
print("-----")
print("Initial Partitions    :", initial_partitions)
print("After Repartitioning  :", updated_partitions)

print("\nExecution Time")
print("-----")
print(f"Before Optimization    : "
      f"{before_optimization:.6f} seconds")

print(f"After Caching        : "
      f"{after_caching:.6f} seconds")

print(f"After Partitioning   : "
      f"{after_partitioning:.6f} seconds")

print("\nPerformance Improvement")
print("-----")
print(f"Caching Improvement    : "
      f"{cache_improvement:.2f}%")

print(f"Partitioning Improvement : "
      f"{partition_improvement:.2f}%")

print("          PERFORMANCE ANALYSIS COMPLETED")
spark.stop()
```

26/09/29 21:32:02 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, 10.255.255.254, 38125, None)

MOVIEFLIX TECHNOLOGIES
SPARK PERFORMANCE OPTIMIZATION

Dataset Loaded Successfully!

INITIAL PARTITION COUNT

Initial Partitions: 1

EXECUTION BEFORE OPTIMIZATION

Average ratings calculated successfully.
Execution Time BEFORE Optimization: 1.392957 seconds

APPLYING CACHE

cache() applied successfully.
Caching triggered using count().
Total Records: 105339
Cache Time: 0.671752 seconds

AVERAGE RATINGS AFTER CACHING

movieId	average_rating
1580	3.627906976744186
471	3.6666666666666665
7253	4.0
3997	2.5
32460	4.5

Run Karo

Run

Python

spark-submit --class org.apache.spark.examples.SparkPi --master spark://localhost:7070

because humne:

loggya hai, normally tumhe Spark ke hundreds of INFO messages nahi dikhne chahiye.

Output approximately aisa dikhega

Tumhare actual numbers obviously different honge.

movieId	average_rating
32460	4.5
144022	3.3125
154190	3.357142857142857
1833	2.95
11591	2.257575757575758
16357	3.6

only showing top 10 rows

Execution Time AFTER Caching: 0.332979 seconds

REPARTITIONING DATASET

Dataset repartitioned successfully.
Updated Partition Count: 4

AVERAGE RATINGS AFTER PARTITIONING

movieId	average_rating
31528	3.5
1342	3.2
4900	4.0
1088	3.419642857142857
2142	3.3181818181818183
4519	3.0454545454545454
26425	4.0
36525	3.5
16357	3.6
16620	3.7916666666666665

only showing top 10 rows

Execution Time AFTER Partitioning: 0.507130 seconds

movieId	average_rating
4519	3.0454545454545454
26425	4.0
36525	3.5
16357	3.6
16620	3.7916666666666665

only showing top 10 rows

Execution Time AFTER Partitioning: 0.507130 seconds

PERFORMANCE REPORT

Dataset : Movie Ratings Dataset
Optimization Techniques : Caching and Partitioning

Partition Information

Initial Partitions : 1
After Repartitioning : 4

Execution Time

Before Optimization : 1.392957 seconds
After Caching : 0.332979 seconds
After Partitioning : 0.507130 seconds

Performance Improvement

Caching Improvement : 76.10%
Partitioning Improvement : 63.59%

PERFORMANCE ANALYSIS COMPLETED

(tushar@Tushar-laptop)~[~/Desktop/BDA_Practical7]

Key Questions /Analysis:

Q1. What is caching in Apache Spark and why does it improve performance?

Caching is a performance optimization technique in Apache Spark. It stores a computed DataFrame or RDD in memory for reuse. When the same data is used again, Spark can use the cached data. It avoids reading and processing the same data repeatedly. Caching is useful when multiple operations are performed on the same dataset. For example, `ratings_df.cache()` is used to cache the ratings DataFrame. The `count()` action can be used to trigger the caching.

Q2. What is the difference between `cache()` and `persist()`?

Both `cache()` and `persist()` are used to store computed data for reuse. `cache()` uses Spark's default storage level. It is simple and convenient to use. `persist()` allows us to explicitly specify the storage level. For example, `df.persist(StorageLevel.DISK_ONLY)` stores data on disk. Therefore, `persist()` provides more flexibility than `cache()`. Both techniques help avoid repeated computation of the same data.

Q3. How does Spark partition data internally?

Spark divides a large dataset into smaller units called partitions. Each partition contains a portion of the complete dataset. Each partition can be processed independently by a Spark task. Multiple partitions can be processed in parallel. This allows Spark to efficiently process large datasets. The number of partitions can be checked using `df.rdd.getNumPartitions()`. Proper partitioning helps improve parallel processing and scalability.

Q4. Why does `repartition()` improve scalability?

`repartition()` redistributes data into a specified number of partitions. For example, `df.repartition(4)` creates four partitions. These partitions can be processed by multiple Spark tasks. In a cluster, the work can also be distributed across multiple executors. This increases parallelism when processing large datasets. Therefore, repartitioning can help Spark applications scale with increasing data. However, `repartition()` causes a shuffle, so it may be slower for small datasets.

Q5. How does lazy evaluation affect Spark performance?

Spark uses a concept called lazy evaluation. Transformations are not executed immediately when they are defined. Instead, Spark builds a logical execution plan for the operations. An action such as `show()` or `count()` triggers the actual execution. Spark can optimize the execution plan before processing the data. This helps avoid unnecessary computations. As a result, lazy evaluation can improve Spark application performance.

Supplementary Problems:

1. Compare `cache()` and `persist(StorageLevel.DISK_ONLY)`. Repartition dataset into 2, 4 and 8 partitions and compare performance. Measure execution time for multiple Spark actions. Use `coalesce()` and compare with `repartition()`. Create a larger ratings dataset and observe scalability.

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import avg
from pyspark import StorageLevel
import time
```

```
# Create Spark Session
spark = SparkSession.builder \
    .appName("MovieFlix_Supplementary") \
    .getOrCreate()

# Hide Spark INFO logs
spark.sparkContext.setLogLevel("ERROR")

print("=====")
print("    MOVIEFLIX TECHNOLOGIES")
print("    SUPPLEMENTARY PROBLEMS")
print("=====")

# Load dataset
ratings_df = spark.read.csv(
    "movie ratings.csv",
    header=True,
    inferSchema=True
)

print("\nDataset Loaded Successfully!")
print("Total Records:", ratings_df.count())

# ----- CACHE -----
cache_df = ratings_df

start_time = time.perf_counter()

cache_df.cache()
cache_df.count()

end_time = time.perf_counter()

cache_time = end_time - start_time

print("\nCACHE()")
print("Storage Level:", cache_df.storageLevel)
print(f"Execution Time: {cache_time:.6f} seconds")

cache_df.unpersist()

# ----- PERSIST DISK ONLY -----
persist_df = ratings_df

start_time = time.perf_counter()

persist_df.persist(StorageLevel.DISK_ONLY)
persist_df.count()
```

```
end_time = time.perf_counter()

persist_time = end_time - start_time

print("\nPERSIST(StorageLevel.DISK_ONLY)")
print("Storage Level:", persist_df.storageLevel)
print(f"Execution Time: {persist_time:.6f} seconds")

persist_df.unpersist()

ratings_df.cache()
ratings_df.count()

# Action 1 - Count
start_time = time.perf_counter()

count_result = ratings_df.count()

end_time = time.perf_counter()

count_time = end_time - start_time

print("\nAction 1 - Count")
print("Records:", count_result)
print(f"Execution Time: {count_time:.6f} seconds")

# Action 2 - Average rating
start_time = time.perf_counter()

average_result = ratings_df.groupBy("movieId") \
    .agg(avg("rating").alias("average_rating"))

average_result.count()

end_time = time.perf_counter()

average_time = end_time - start_time

print("\nAction 2 - Average Rating")
print(f"Execution Time: {average_time:.6f} seconds")

# Action 3 - Filter
start_time = time.perf_counter()

high_rating_count = ratings_df.filter(
    ratings_df.rating >= 4.0
).count()

end_time = time.perf_counter()
```

```
filter_time = end_time - start_time

print("\nAction 3 - Filter Rating >= 4.0")
print("Records:", high_rating_count)
print(f"Execution Time: {filter_time:.6f} seconds")

ratings_df.unpersist()

# Create 8 partitions first
base_df = ratings_df.repartition(8)

print("\nInitial Partitions:", base_df.rdd.getNumPartitions())

# ----- COALESCE -----
start_time = time.perf_counter()

coalesce_df = base_df.coalesce(2)

coalesce_df.count()

end_time = time.perf_counter()

coalesce_time = end_time - start_time

print("\nCOALESCE(2)")
print("Partitions:", coalesce_df.rdd.getNumPartitions())
print(f"Execution Time: {coalesce_time:.6f} seconds")

# ----- REPARTITION -----
start_time = time.perf_counter()

repartition_df = base_df.repartition(2)

repartition_df.count()

end_time = time.perf_counter()

repartition_time = end_time - start_time

print("\nREPARTITION(2)")
print("Partitions:", repartition_df.rdd.getNumPartitions())
print(f"Execution Time: {repartition_time:.6f} seconds")

original_count = ratings_df.count()

# Create a larger dataset by repeating the original dataset
large_df = ratings_df
```

```
for i in range(4):
    large_df = large_df.unionByName(ratings_df)

start_time = time.perf_counter()

large_count = large_df.count()

end_time = time.perf_counter()

large_count_time = end_time - start_time

print("\nOriginal Dataset")
print("Records:", original_count)

print("\nLarger Dataset")
print("Records:", large_count)
print(f"Execution Time: {large_count_time:.6f} seconds")

# Average rating on larger dataset
start_time = time.perf_counter()

large_average = large_df.groupBy("movieId") \
    .agg(avg("rating").alias("average_rating"))

large_average.count()

end_time = time.perf_counter()

large_average_time = end_time - start_time

print(f"Average Rating Execution Time: {large_average_time:.6f} seconds")

print("\nCache vs Persist")
print("-----")
print(f"cache() Time : {cache_time:.6f} seconds")
print(f"persist(DISK_ONLY) Time : {persist_time:.6f} seconds")

print("\nMultiple Spark Actions")
print("-----")
print(f"Count Action : {count_time:.6f} seconds")
print(f"Average Rating Action : {average_time:.6f} seconds")
print(f"Filter Action : {filter_time:.6f} seconds")

print("\nCoalesce vs Repartition")
print("-----")
print(f"coalesce(2) Time : {coalesce_time:.6f} seconds")
print(f"repartition(2) Time : {repartition_time:.6f} seconds")

print("\nScalability")
print("-----")
```

```

print(f"Original Records      : {original_count}")
print(f"Larger Dataset Records  : {large_count}")
print(f"Larger Dataset Time     : {large_count_time:.6f} seconds")
print(f"Larger Average Time     : {large_average_time:.6f} seconds")
spark.stop()

```

The screenshot displays a Jupyter Notebook interface with a dark theme. The left sidebar shows the notebook's structure, including sections for 'MOVIEFLIX TECHNOLOGIES', 'SUPPLEMENTARY PROBLEMS', and 'Dataset Loaded Successfully!'. The main area shows the execution of Spark code, with output logs for 'CACHE() VS PERSIST(DISK_ONLY)', 'EXECUTION TIME FOR MULTIPLE SPARK ACTIONS', and 'COALESCE() VS REPARTITION()'. A 'SUPPLEMENTARY SUMMARY' table is also visible, summarizing the execution times for various operations.

SUPPLEMENTARY SUMMARY	
Cache vs Persist	
cache() Time	: 0.843460 seconds
persist(DISK_ONLY) Time	: 0.483932 seconds
Multiple Spark Actions	
Count Action	: 0.128303 seconds
Average Rating Action	: 0.728351 seconds
Filter Action	: 0.274917 seconds
Coalesce vs Repartition	
coalesce(2) Time	: 0.232448 seconds
repartition(2) Time	: 0.164122 seconds
Scalability	
Original Records	: 105339
Larger Dataset Records	: 526695
Larger Dataset Time	: 0.238380 seconds
Larger Average Time	: 0.595538 seconds

Tools / Technology used:

Apache Spark 3.x | PySpark | Python 3.11 | VS Code | Jupyter Notebook

Conclusion

The Practical successfully demonstrated how Apache Spark caching and partitioning can be used to optimize the performance of data processing applications. The movie ratings dataset was loaded and analyzed using Spark DataFrames, and the execution time was measured before and after applying optimization techniques. Caching helped reuse the computed dataset without repeated processing, while repartitioning distributed the data into multiple partitions for parallel processing. The performance of different partition counts such as 2, 4, and 8 was also compared. The practical showed that the effect of these optimization

techniques depends on the dataset size, workload, and number of partitions. Overall, Spark caching and partitioning are useful techniques for improving processing efficiency, scalability, and performance in Big Data applications.

Github Repo:

<https://github.com/tushar-2606/BDA-SHP-2026>

